DOUBT SESSION

- Design of FB Messenger
  - a.) How to choose DB
- Bloom filters
  - a.) UU Beautiful Algo
- Doubts

45 mins

This Saturday we won't be taking class.

Design of FB Messenger

Choosing DB

A : W ⇒ 1 : 1

⇒ Both read as well as write heavy

⇒ Try to reduce one of them

Let's try to reduce write

eg. Google Typeahead.

→ Batched write

✗ ⇒ This won't work in FB Messenger

⇒ Consistency is very very important here.

Let's try to reduce reads

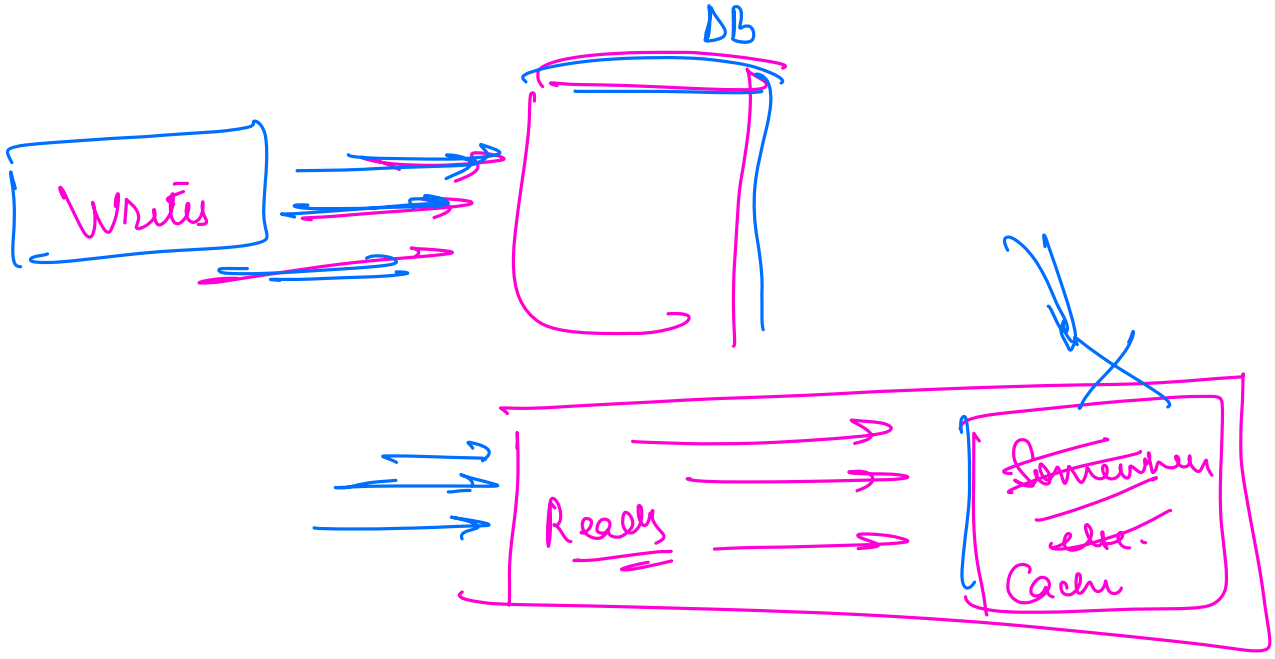

⇒ Because I have to maintain consistency all data in DB should also be present in 'X'

→ X Should be consistent with DB

→ X Should be able to handle large # of write

⇒ X should be able to handle large # of reads.

⇒ X should be data other than disk

⇒ X should be serving data in RAM.

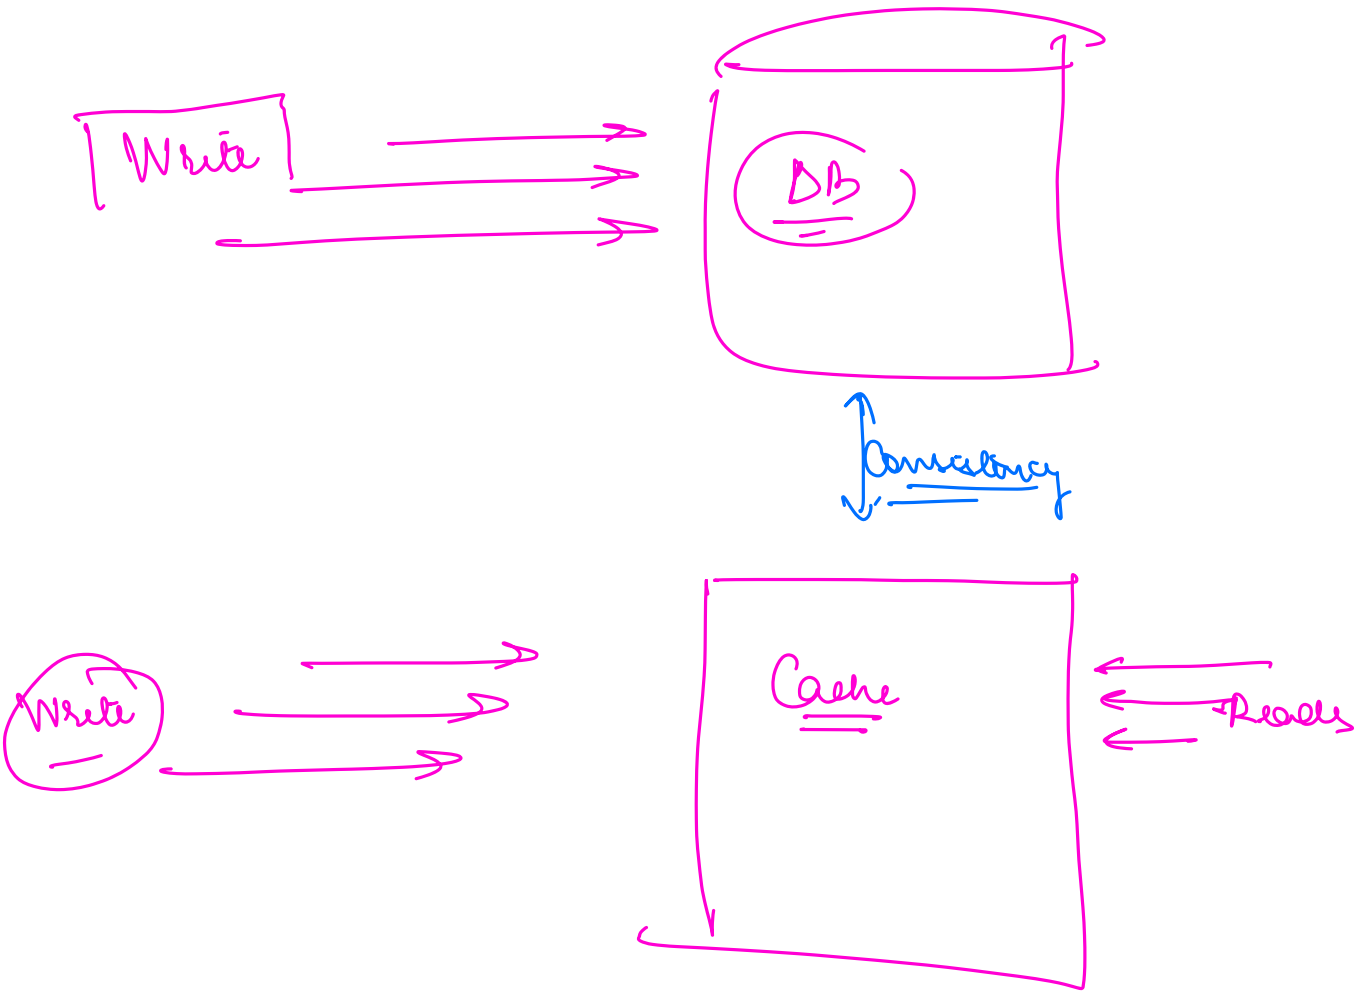

⇒ Write Through Cache

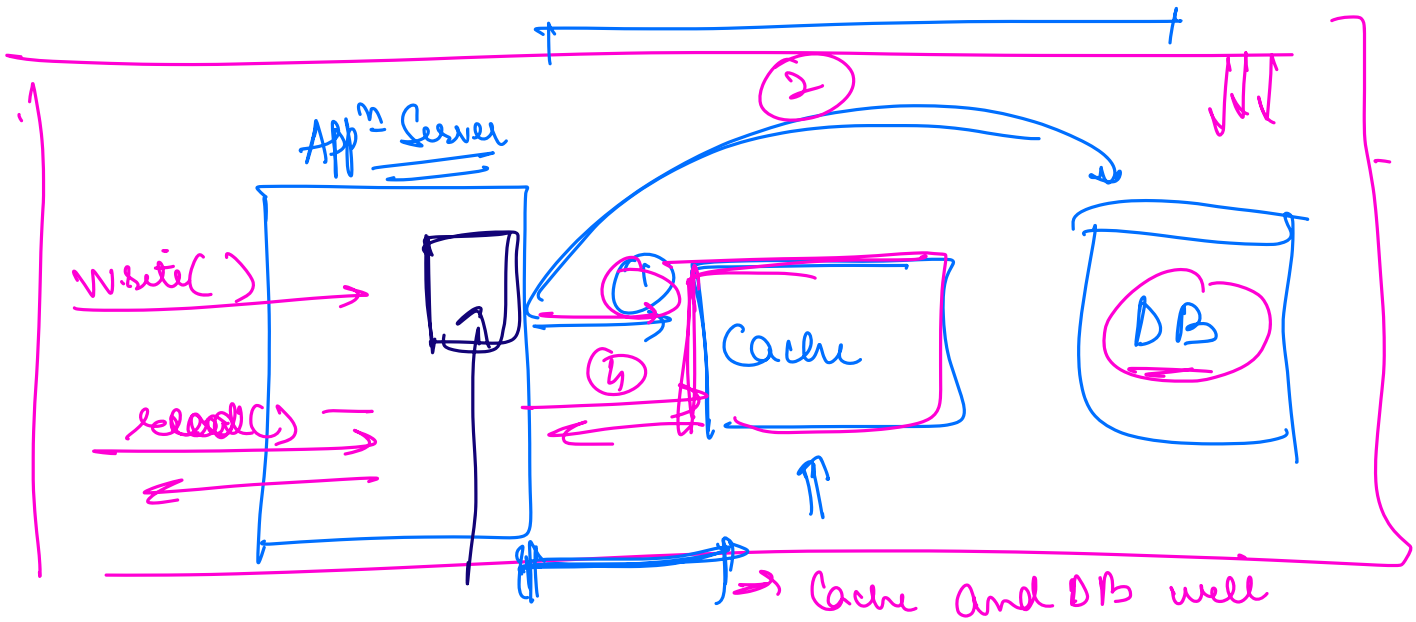

→ Cache and DB will always be consistent

Can I optimize this further

Latency while writing

(i) app^n Server and cache

(ii) app^n Server and DB

⇒ Slow system

Latency while reading

(i) app^n Server and Cache

What if I cache within the app^n Server

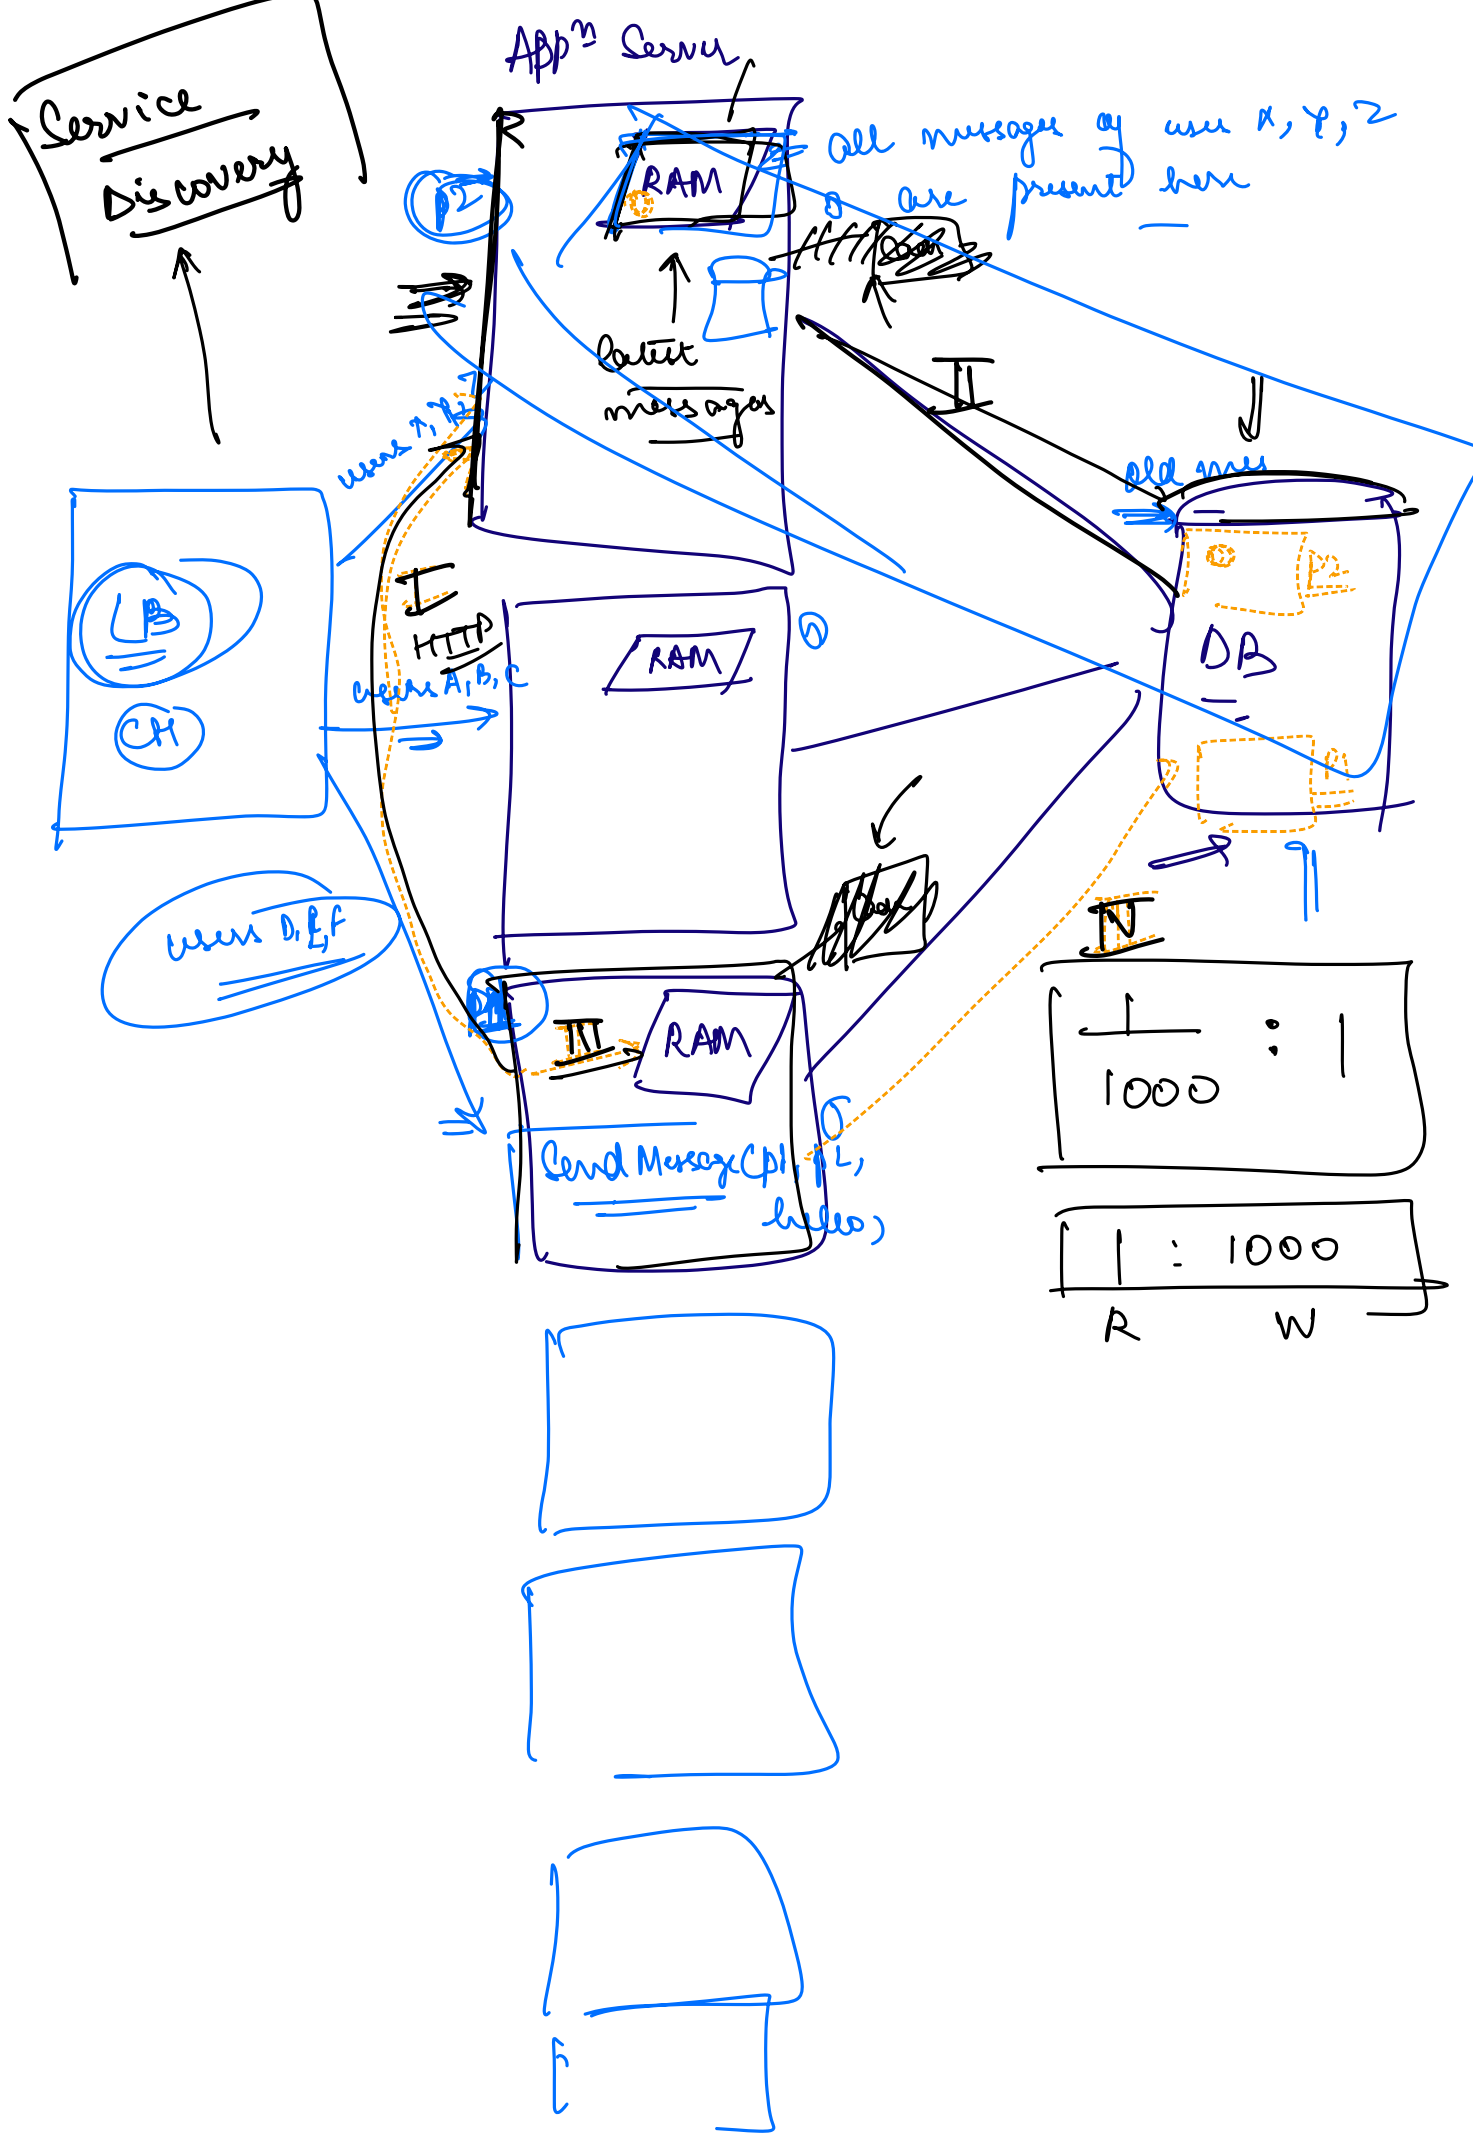
Service Discovery
App^n Servers
RAM
all messages of users X, Y, Z are present here
Collect messages
users X, Y, Z
old users
LB
CH
HTTP
users A, B, C
RAM
DB
users D, E, F
RAM
Send Message (pk, pL, hello)
1 : 1000
1000 : 1
R
W

Send Message (P1) (P2) hello) {

⇒ { we have to put data in shards of P1 and P2 }

}

What DB ?

- ⇒ Handle large # of write
- ⇒ NoSQL DB like Cassandra / HBase
  - ⇒ |R|, |W|

small value of W

Document DB ? unstructured data

message

| id | sent by | time | — | — | — |

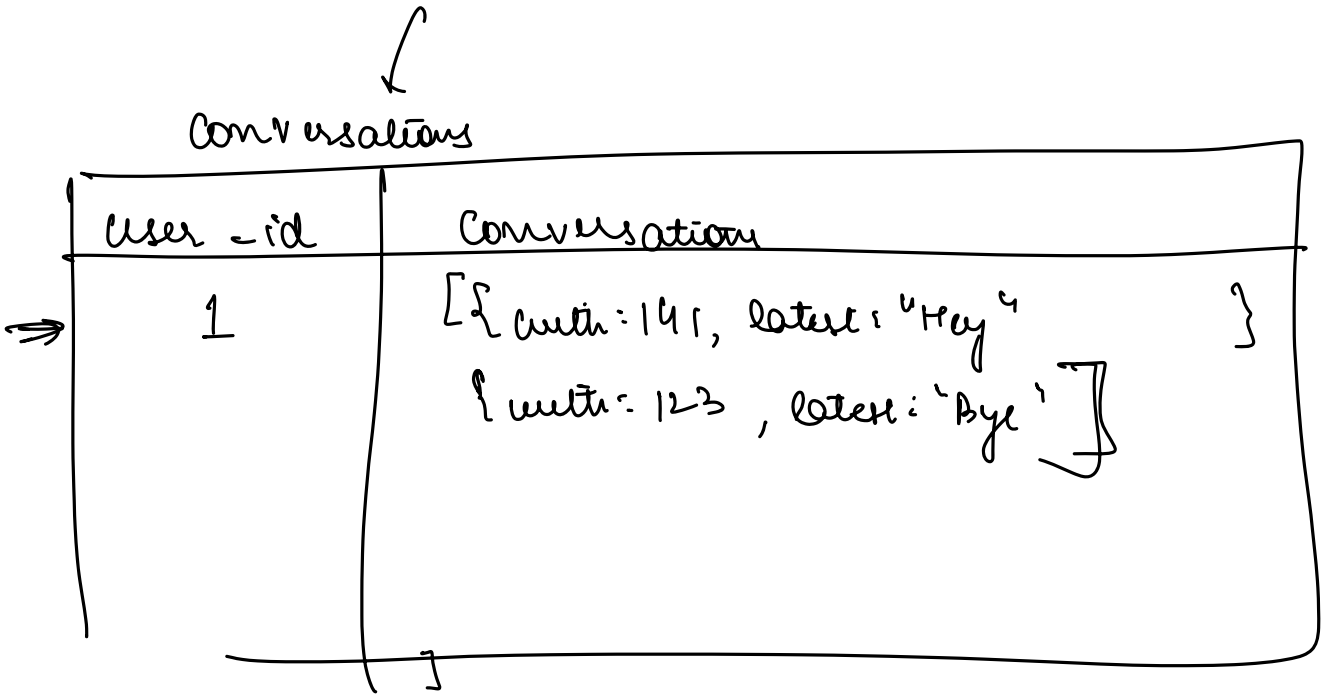

Bloom filters

- → to search for something you have to go to DB
- → to check if something exists or not, we need to go to DB.

⇒ latency

⇒ time DB will take to execute query

NoSQL Internals

⇒ LSM Trees

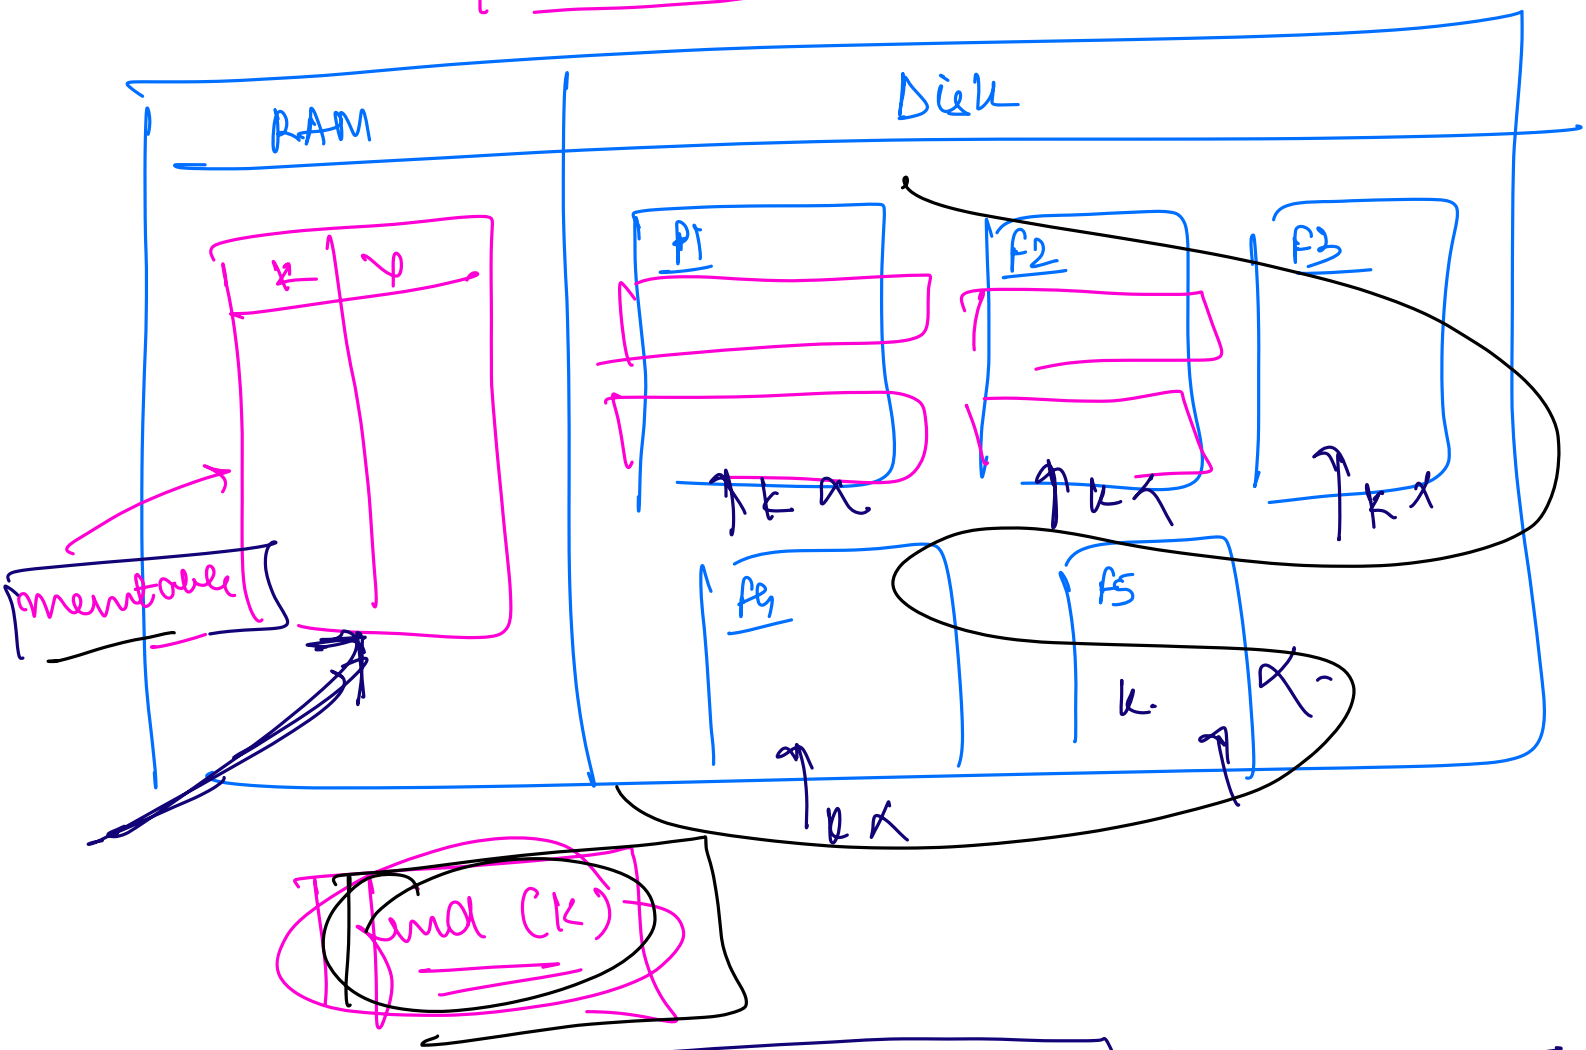

⇒

Go to memtable: if yes: return

if no: find one by one in each file from latest to oldest.

if still not there ⇒ doesn't exist

⇒ { Searching key that doesn't exist is going to be most time consuming }

{ "I have to spend a lot of time to check if something doesn't exist" }

Can I get to know of something doesn't exist fast ⇒

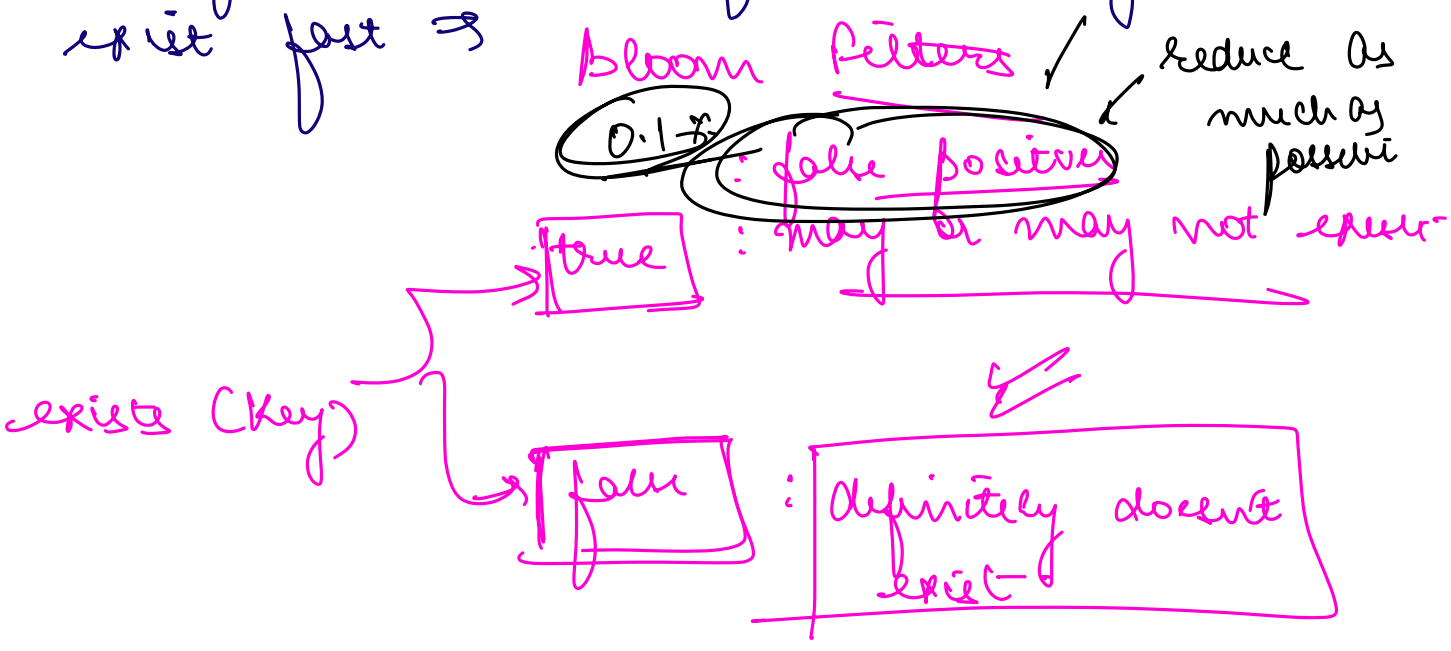

⇒ 99.9% of times if a key doesn't exist ⇒ done

⇒ 0.1% chance where I will still check

How Bloom filters Work

① Bit Array → initially everything 0

② Bigger a bloom filter, lower false positive
The more the hash fn, lower false positive

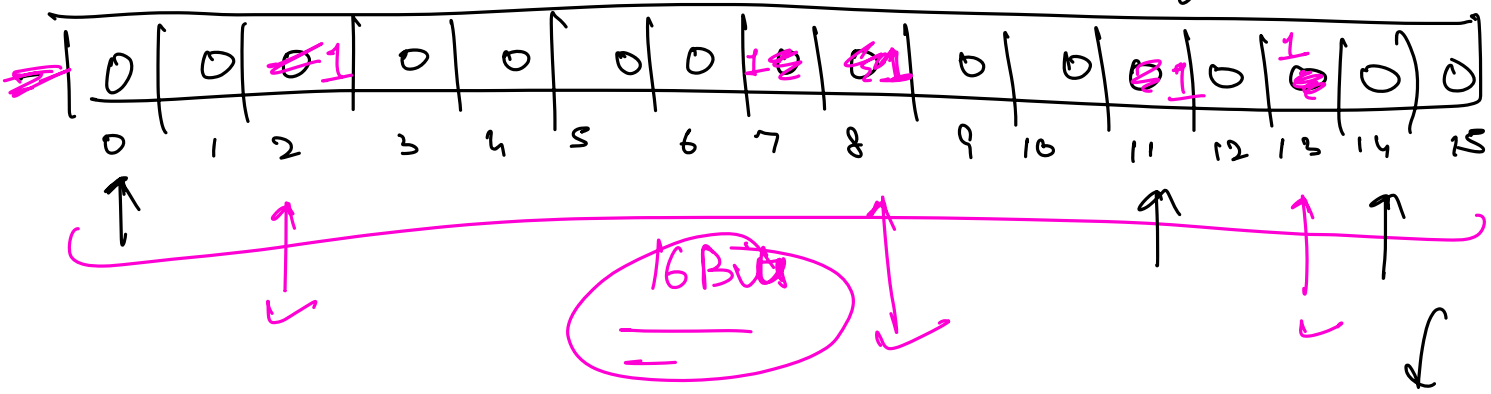

③ Multiple hash fn, each selects from 0 to 15

- h1( )
- h2( )
- h3( )

Whenever someone enters a key :

Compute the hash value of that key for all hashes.

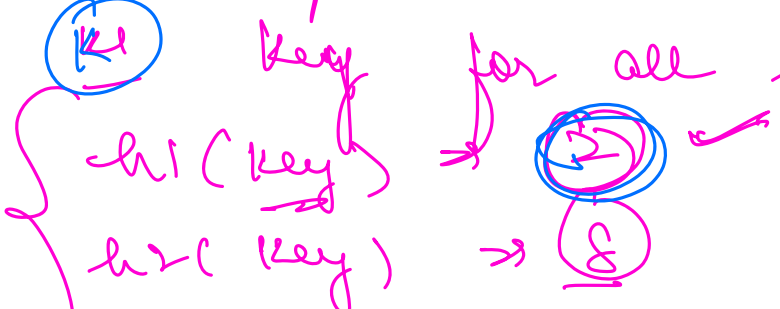

h3(key) ⇒ #1

Set (1) in hit array at all three places.

Insert (k2)

h1(k2) → 8
h2(k2) → 7
h3(k2) → 13

} O(1)

When someone searches a key

Search (key)

⇒ Compute value of all hash f^{n} for that key

Search (k3) → h1(k3) ⇒ 0
h2(k3) ⇒ 11
h3(k3) ⇒ 14

0 → 2
11 → 13
14 → 8

⇒ Check if in bitarray all of these are

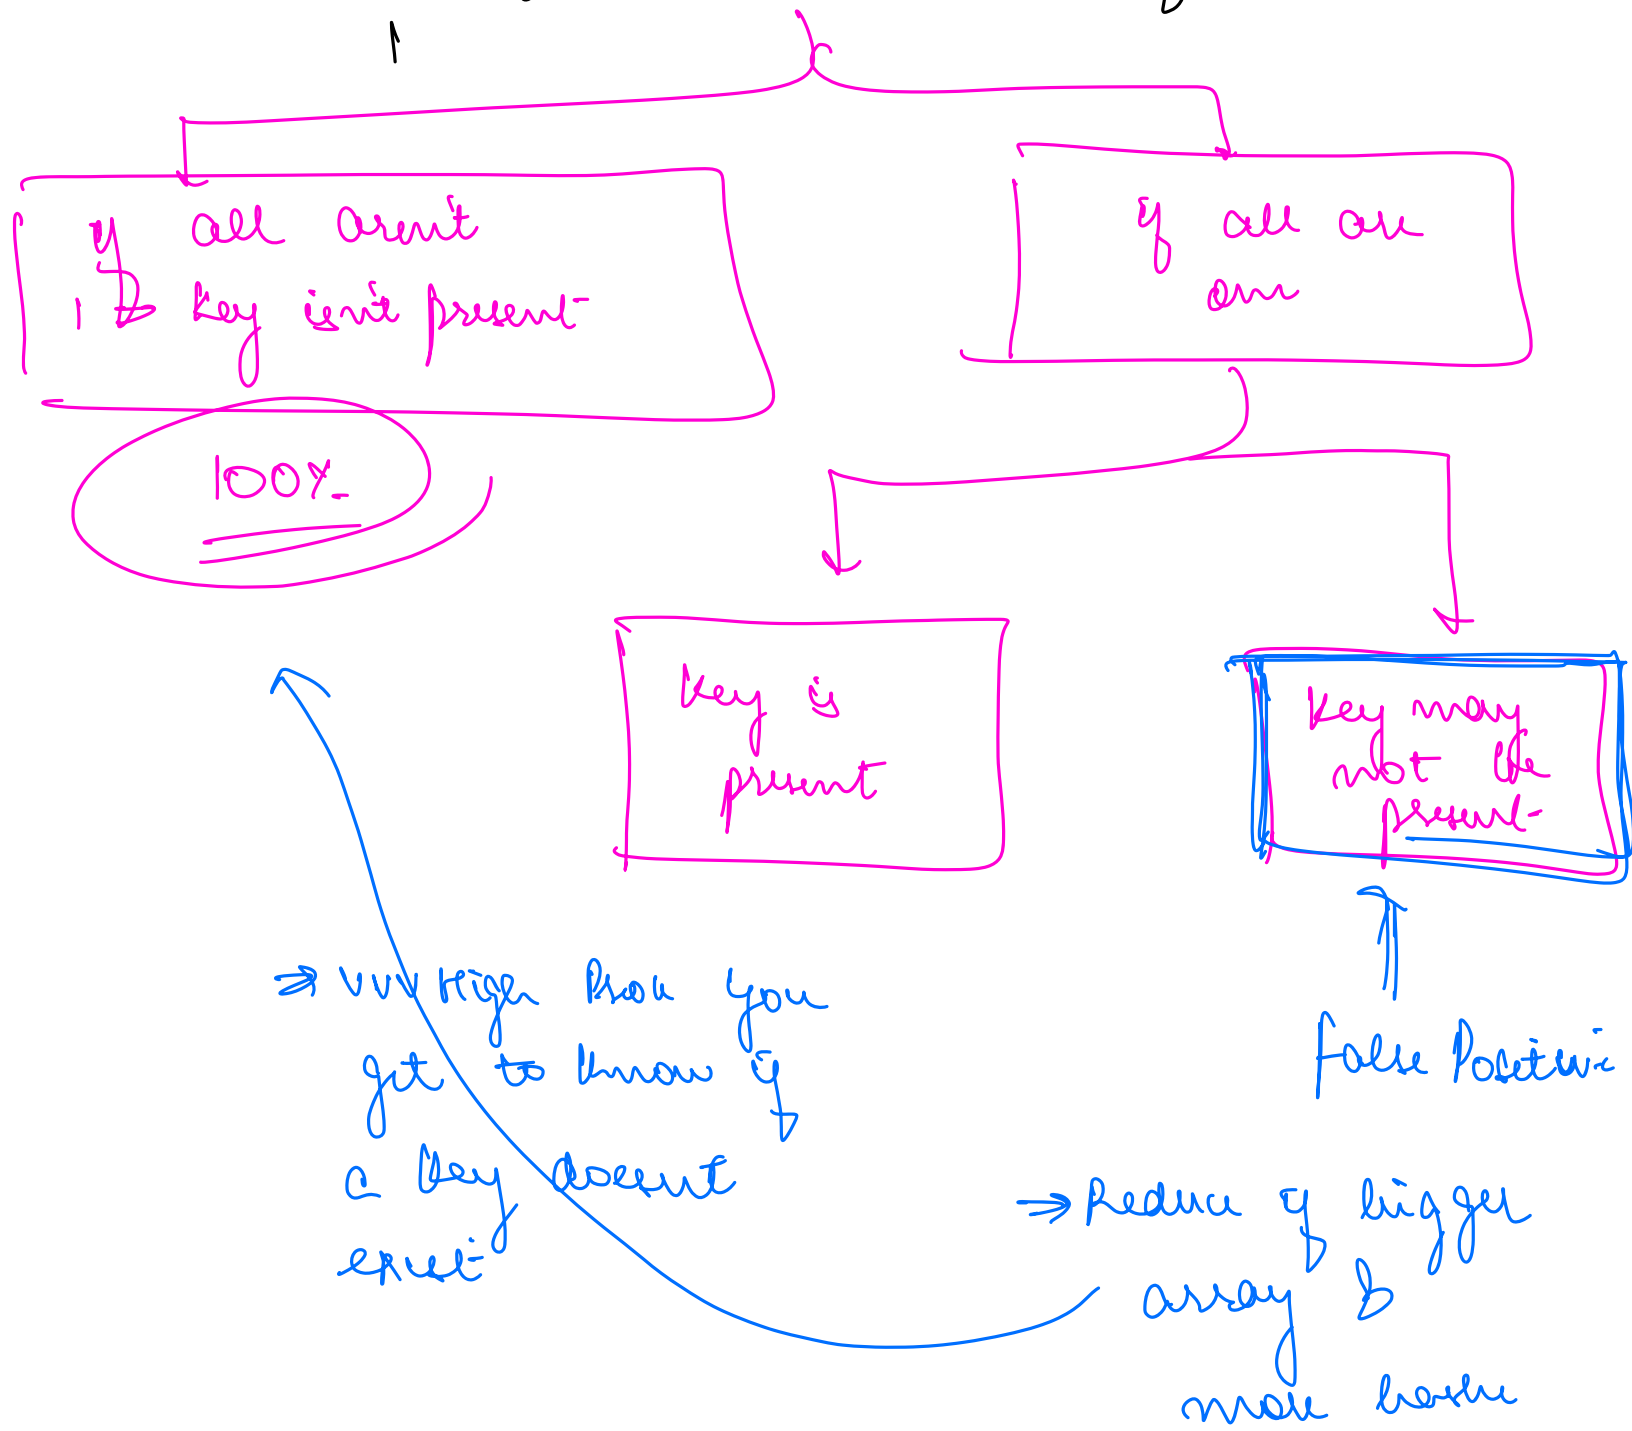

⇒ VVV High Risk you get to know if a key doesn't exist

⇒ Reduce if bigger array & more hashes

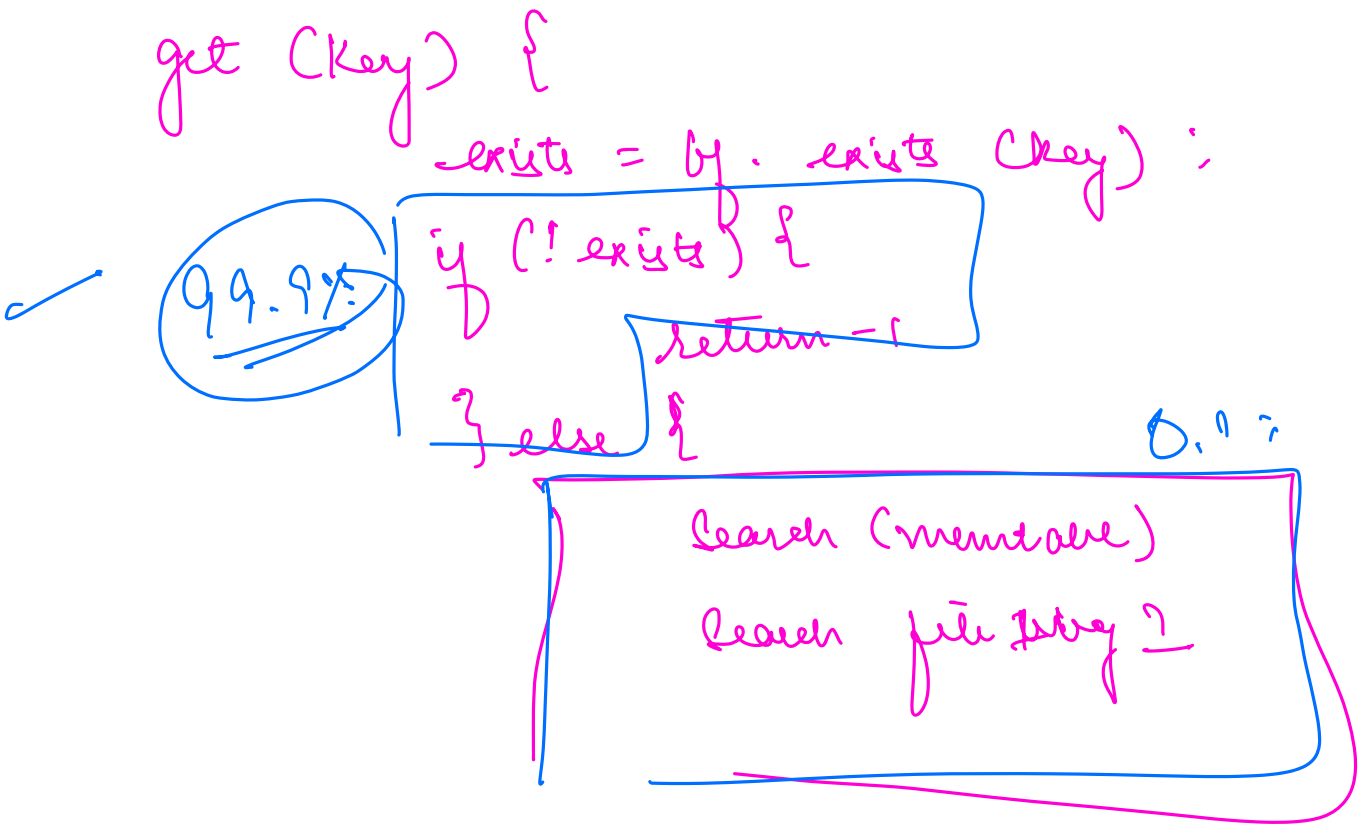

}

insert (key) {

bf. insert (key) → O(1)

}

[100000] By [2 day]

1 billion → 32 bits

10 Million

^ 10 Million

[dbr.put(1, { })] =

Cassandra

delete (key) → don't do anything

⇒ every day reconstruct

delete (k4)

h1(k4) → 4 → 0

h2(k4) → 2 → 0

h3(k4) → 8 → 0

2×10^9 ops / sec

$$\frac{50}{2 * 10^9} \text{ sec}$$